

The claim intake agent that refuses to guess

Voice claim intake on the AssemblyAI Voice Agent API

It cannot write a value into the record unless a server-side validator approves it.

`github.com/sadishihab/claim-intake-agent`
`web-production-74b77.up.railway.app`

The problem

ONE POLICY NUMBER, READ ALOUD

3841188

M-A-C-K-K-K-D-4-1-1-3-8

D411

KD41118

KD41282

All rejected. None is a policy on file.

THE HARDER CASE

BX7 - 4402 Marcus Halloway

BX7 - 4420 Priya Raghunathan

One transposition apart. Two different people.

**A wrong value that looks right
gets filed and never surfaces.**

The architecture



The model proposes a value. It never writes one.

The verdict carries the exact words to speak, and the prompt requires them verbatim.
Audio is proxied through the server, so the API key never reaches the browser.

The three verdicts

accepted

checked and unambiguous

records it, says nothing, moves on

unconfirmed

probably right, cannot be sure

reads the value back and waits

rejected

cannot be right

explains the problem and asks again

The middle one is where the design went.

A system with only accept and reject guesses on every near miss.

BX7-4402

"four four zero two"

BX7-4420

"four four two zero"

Two real policies. The readback separates them.

What we measured

Three null results on recogniser tuning.

Baseline, 12 clips	12 / 12	Test at its ceiling. Cannot discriminate.
Noise at 15, 10 and 5 dB SNR	no change	Still transcribed perfectly.
transcription_prompt, max_accuracy	no effect	Not carried.

The clips were text-to-speech, which cannot produce human disfluency.

The harness was a scratch script and is not in the repository.

One lever did move something. It made the system worse.

The keyterms finding

input.keyterms biased transcription toward the policy numbers in policies.json.

```
partial    Yes, my policy number is C411.
partial    Yes, my policy number is KD4-1188.          <- revised onto a keyterm
FINAL      Yes, my policy number is KD4-1188.
recorded    KD4-1188    accepted    exact match on file
```

**Never bias the recogniser toward the answer key
when the recogniser's output is the evidence being validated.**

Every downstream guard still passed. Exact match was only evidence while the recogniser knew nothing about the policy list.

The audio is not kept. That the number read was wrong is my recollection, not something the log proves.

The claim intake agent that refuses to guess

The evidence trail

Every attempt is appended to calls/<session_id>.jsonl as it happens.

25.0s	policy_number	KD41188	UNCONFIRMED	
	heard	"My policy number is KD41188."		
	readback	"That's Kilo Delta four, one one eight eight, correct?"		

41.5s	policy_number	KD41188	ACCEPTED	-> KD4-1188
	heard	"Yes, correct."		

A policy number is never accepted on a match alone.

The caller has to hear it read back and agree. ClaimRecord enforces that, not the prompt.

/compare renders six moments side by side. The panel at / tails the same log. 196 tests.

The claim intake agent that refuses to guess